

Anton Dziatkovskii

Developer activation for frontier platforms · operating lead, Palo Alto AI Research Lab (an independent lab)
Email: dzyatkovskiy.a@gmail.com · WhatsApp: +1 341 222 9178 · Calendly:
<https://calendly.com/paloaltolab/1-on-1> · Web: <https://tonydz.github.io/> · Palo Alto, CA / Sintra, PT · US
O-1 visa (active) · EU citizen (Polish passport) — no visa sponsorship required

SUMMARY

I've done developer activation for a frontier platform before — when smart contracts were where LLM agents are now. Classically trained in CS security (MEPhI), programming in Python and C++, eleven years as a hired operating executive (COO/CTO) at a startup incubator and software house (in crypto since 2015 — supported early Ethereum and Cardano fundraising efforts in Japan), author of two preprints on multi-agent stability. A superconnector with a builder's toolkit: a decade-deep network of founders, builders, engineers and investors (10,000+ founder calls, 2,000+ institutional-investor calls) — and I automate my own connector work end-to-end, from a self-built CRM to today's AI-agent fleet. Today I run a production multi-machine Claude fleet in public — and write the operating manual as I go. Built with Claude as implementation collaborator; I own problem framing, architecture, evaluation, product decisions and final QA. Based between Palo Alto, CA — the heart of Silicon Valley — and Sintra, PT. US O-1 visa active and EU citizen (Polish passport): authorized to work in the US and anywhere in the EU immediately, no visa sponsorship required.

SKILLS

Python, C++, CRM, Salesforce, Microsoft Navision, AI agents, Model Context Protocol (MCP), RAG, evals, developer activation, fundraising, partnerships, Solidity

EXPERIENCE

Operating Lead — Palo Alto AI Research Lab (2023 – present, Palo Alto, CA)

Grew a community of 100+ AI practitioners, including Stanford researchers and big-tech engineers. Production multi-machine Claude fleet run in public: claw-consensus (reproducible multi-machine agent consensus, offline demo, published evals, FAILURE-MODES.md), verbatim-citation-gate (zero-token gate + burden-of-proof judge that catches fabricated RAG citations, MIT), agent-control-plane-casebook (reproducible control-plane failures from the production fleet — each case ships a deterministic repro, a runnable red test and an upstream bug report), the operating manual written day by day (相棒 AIBŌ · The Partner).

- Red-team my own production agent fleet and publish the result: a written threat model (8 assets, 6 trust boundaries, 10 untrusted input surfaces, 9 residual risks carrying numbers) mapped onto OWASP Top 10 for LLM Applications (2025) and MITRE ATLAS by exact identifier — published together with its holes: 15 ATLAS techniques reproduced on my own stand, 10 more named as applicable and untested, 12 ATLAS mitigations implemented and 8 explicitly not ([agent-fleet-red-team](#))
- Measured the fleet's human-approval gate as two separate numbers rather than one accuracy figure (2026-10-08): 100% refusal across 49 must-ask actions, 50% over-refusal across 42 benign ones, and 71.4% rule coverage — 14 of the 49 reach a human only because the default is fail-closed, each one named in the report. Six-way mutation suite, shown failing on broken code before it was shown passing. The same gate on real traffic: 2,879 high-risk firings in 90 days, 0 asks expired unanswered
- Reproduced an indirect prompt-injection chain end to end on my own stand ([leash-poc](#)): one poisoned string, four input surfaces, four agents, 24 requests. Naive fleet exfiltrated 168 bytes, deleted 5 files, rotated 4 production keys and spawned 4 admin sub-agents; behind a single independent authorization

boundary, zero on all five — with the user's work still completed. 27 hash-chained decisions, chain verified after SIGKILL mid-run

- Merged into microsoft/semantic-kernel (#14371, Oct 2026): closed the check-time/use-time gap in the OpenAPI plugin's anti-SSRF validator. The request now goes to the exact address the validator vetted, so on the plugin's built-in HTTP client a DNS-rebinding hostname can no longer pass the check and then reach an internal address (scoped as hardening; caller-supplied clients and proxies stay out of scope); TLS is still verified against the hostname, no new dependency, 12 tests including one that drives real httpx + httplib through a simulated rebind
- Run an evidence-first reliability program across the AI-agent open-source ecosystem: reproducible cross-project probes, mutation-tested reviews of other people's pull requests, and bug reports that arrive with a deterministic reproduction and a regression test demonstrated failing on the unfixed code
- Review accepted by the author, not merely posted — google/adk-python#6957, an OAuth2 client secret leaking over three endpoints: the mutation matrix showed which of the project's own tests stayed green, and the author adopted all three findings within six hours ("genuinely one of the most useful reviews I've gotten on this PR")
- Failure classes driven to an upstream fix: google/adk-python#6887 (session ids normalised on write but not on read) fixed from the report; #6941, filed with a measured root cause, was labelled core, assigned to a maintainer and answered by two third-party pull requests
- Credited in a third-party release 40 hours after the report — headroom v2.0.8: "reported by @tonydzi"
- Scope of the program since Jul 2026 (verified 2026-10-08): 40 pull requests merged into 25 third-party engineering projects — Microsoft Semantic Kernel, Model Context Protocol Go SDK, UK AISI inspect_ai, QwenLM/qwen-code, google-gemini/cookbook, bytedance/deer-flow, fastmcp, agno, pydantic/logfire, deepset-ai/haystack-integrations — plus 65 issues filed upstream
- The program runs unattended: 13 daily autonomous lanes on the fleet (code review, live-install documentation QA, issue-first bug reports, thread watch), with human approval gates on anything irreversible

COO / CTO (hired executive) — Platinum Software Development Company / Platinum VC & Incubator (2015 – present, APAC)

Hired operating executive (COO/CTO) who ran a startup incubator + software house in APAC since 2015 (11 years), alongside Japanese partner Tetsuji Nagata (Grand Corp; 12 years at Bloomberg before that). Deal flow, cohorts, founder coaching, advisor matching, fundraising — hands-on, with rare depth in Silicon Valley and even the Japanese and Korean startup communities. Portfolio ~70% crypto / 30% other, AI startups included.

- Superconnector as a system: founders, builders, engineers and investors matched daily — every repeatable connector motion (enrichment, matching, follow-up) automated, first with a self-built CRM, today with AI agents
- In fundraising since 2017: \$35M fund under management, syndicate of 250 VCs/angels; helped portfolio startups raise end-to-end — deck, investor intros, term negotiation
- Built a deal-flow machine: 10,000+ founder calls and 2,000+ institutional-investor calls; wrote my own CRM to run it (up to 30 operators daily; today it is AI-native)
- Shipped crypto-exchange infrastructure: cryptocurrency backends, market-maker engines, high-frequency trading systems; partnerships with hedge funds, market makers, brokers
- In 2018 worked with banks and governments in Myanmar, Cambodia, Thailand and Australia (incl. CB Bank in Myanmar and central-bank stakeholders) on frameworks to launch their first stablecoins
- Built custom distributed ledgers for enterprise clients (ANA, Foxconn) before "blockchain" was an industry term; enterprise advisory for subsidiaries of Foxconn and ANA Airlines

- Built Solidity curricula and ran hackathons, cohorts and incubations for Web3 builders — from first workshop to shipped mainnet product; scaled a distributed engineering org to 40+ developers as product owner
- Worked with major blockchain foundations — Ethereum, Avalanche, Cardano, Aptos, Sui, Canton — supplying engineers and deal flow ("engineers are the new oil")

Chief Engineer, Lending & Risk Management — Everex (Singapore) (2017 – 2018)

Wrote smart contracts and led government/bank relations across SE Asia on structuring crypto payments and stablecoins; blockchain credit scoring, stable-value payments.

Sales Director, compute hardware (Southeast Europe) / CRM-ERP Development Director — Merlion (2006 – 2015)

One of the largest private IT distributors in Eastern Europe. Sold compute hardware to B2B and enterprise buyers: semiconductors and chips, memory, GPUs, CPUs and the rest of the computing stack, with Southeast Europe as my sales territory. In parallel built the ERP and CRM systems the sales floor ran on (Microsoft Navision and Salesforce, 5 years), where I learned that deal flow runs on CRM quality.

- Carried Southeast Europe as sales territory for compute hardware: chips and semiconductors, memory, GPUs, CPUs and adjacent computing equipment, sold B2B and enterprise to customers procuring compute capacity and hardware infrastructure
- Personal turnover ~\$5M per month; the department I sold in ran \$15M to \$25M per month
- Sold on a deliberately competitive floor: departments competed against each other for the same enterprise accounts

CRM Developer Intern — Salesforce/ERP ecosystem (2004 – 2006)

First professional years: CRM systems development in the Salesforce/ERP ecosystem.

EDUCATION

PhD in Education (Information Technologies) — Paris College of International Education, 2021
 MSc, Computer and Information Systems Security — National Research Nuclear University MEPhI
 BSc, Computer and Information Sciences — National Research Nuclear University MEPhI

SELECTED PUBLICATIONS & RESEARCH

Homeostatic Governance for multi-agent systems (three-circuit control) — Preprint, arXiv submission in progress.

Le Chatelier Tokenomics (multi-asset coupled equilibria) — Preprint, arXiv submission in progress.

50+ published items, 2018–2026 (137 citations, h-index 7, verified 2026-09-01) —

<https://tonydzi.github.io/scholar/publications/>

Full academic profile: <https://tonydzi.github.io/scholar/>

LANGUAGES

English (Working proficiency) · Russian (Native) · Ukrainian (Working proficiency) · Polish (Working proficiency)

INTERESTS

Endurance sport: triathlon, mid-pack drifting to the back and honest about it; half-marathons; surfing, windsurfing, kitesurfing, tennis; mountain biking (downhill, jumps) on an analogue bike by choice, because a workout should stay a workout.

Aviation and building bunkers: licenced helicopter pilot, ~200 flight hours; student pilot on fixed-wing; I build bunkers - several already standing in New Zealand and Australia, and the fixed-wing licence is how I

reach them when things go sideways; long-term goal: a small plane for the family, and a round-the-world flight in it.

PROFILES

GitHub: <https://github.com/tonydzi> | X: https://x.com/Tony_Stef_ | Google Scholar:
<https://scholar.google.com/citations?user=b8gKHiMAAAAJ> | ORCID: <https://orcid.org/0000-0001-7408-3054> | Academia.edu: <https://tylip.academia.edu/AntonDziatkovskii> | Wellfound:
<https://wellfound.com/u/anton-dziatkovskii-3>